

We start by installing the MDP library that we will use throughout the workshop.

```
1 !pip install mdp
```

```
Requirement already satisfied: mdp in /usr/local/lib/python3.7/dist-packages (
Requirement already satisfied: future in /usr/local/lib/python3.7/dist-package
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages
```

```
1 import os, inspect
2 stdlib = os.path.dirname(inspect.getfile(os))
3 print(stdlib)
```

```
/usr/lib/python3.7
```

And we now mount our Drive to make use of the additional files. Here I have put these on the Colab Notebooks directory on my Drive. If you use a different directory, the path should be changed accordingly.

```
1 from google.colab import drive
2 import sys
3 drive.mount('/content/drive')
4 sys.path.insert(0, '/content/drive/MyDrive/Colab Notebooks')
5 import _mdp as mdp
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call

We will now setup the practical script and the virtual display for the visualisation tasks throughout. The cell may have to be run a couple of times so as to allow xvfb to start.

```
1 import sys, os
2 if 'google.colab' in sys.modules and not os.path.exists('.setup_complete'):
3     # Insert the directoryimport sys
4     # Run the scripts
5     !setup_colab_practical3.sh -O- | bash
6     !touch .setup_complete
7
8 # This code creates a virtual display to draw game images on.
9 # It will have no effect if your machine has a monitor.
10 if type(os.environ.get("DISPLAY")) is not str or len(os.environ.get("DISPLAY"))
11     !bash ../xvfb start
12     os.environ['DISPLAY'] = ':1'
```

```
1 transition_probs = {
2     's0': {
3         'a0': {'s0': 0.5, 's2': 0.5},
4         'a1': {'s2': 1}
5     },
6     's1': {
```

```

7         'a0': {'s0': 0.7, 's1': 0.1, 's2': 0.2},
8         'a1': {'s1': 0.95, 's2': 0.05}
9     },
10    's2': {
11        'a0': {'s0': 0.4, 's2': 0.6},
12        'a1': {'s0': 0.3, 's1': 0.3, 's2': 0.4}
13    }
14 }
15 rewards = {
16     's1': {'a0': {'s0': +5}},
17     's2': {'a1': {'s0': -1}}
18 }
19 from _mdp import MDP
20 mdp = MDP(transition_probs, rewards, initial_state='s0')

```

We can now use MDP just as any other gym environment:

```

1  print('initial state =', mdp.reset())
2  next_state, reward, done, info = mdp.step('a1')
3  print('next_state = %s, reward = %s, done = %s' % (next_state, reward, done))

initial state = s0
next_state = s2, reward = 0.0, done = False

```

but it also has other methods that you'll need for Value Iteration

```

1  print("mdp.get_all_states =", mdp.get_all_states())
2  print("mdp.get_possible_actions('s1') = ", mdp.get_possible_actions('s1'))
3  print("mdp.get_next_states('s1', 'a0') = ", mdp.get_next_states('s1', 'a0'))
4  print("mdp.get_reward('s1', 'a0', 's0') = ", mdp.get_reward('s1', 'a0', 's0'))
5  print("mdp.get_transition_prob('s1', 'a0', 's0') = ", mdp.get_transition_prob(

mdp.get_all_states = ('s0', 's1', 's2')
mdp.get_possible_actions('s1') = ('a0', 'a1')
mdp.get_next_states('s1', 'a0') = {'s0': 0.7, 's1': 0.1, 's2': 0.2}
mdp.get_reward('s1', 'a0', 's0') = 5
mdp.get_transition_prob('s1', 'a0', 's0') = 0.7

```

## ▼ Visualising MDPs

You can visualize any MDP with the drawing fuction donated by [neer201](#). To do so, we have to install graphviz.

```

1  from _mdp import has_graphviz
2  from IPython.display import display
3  print("Graphviz available:", has_graphviz)

Graphviz available: True

```

```

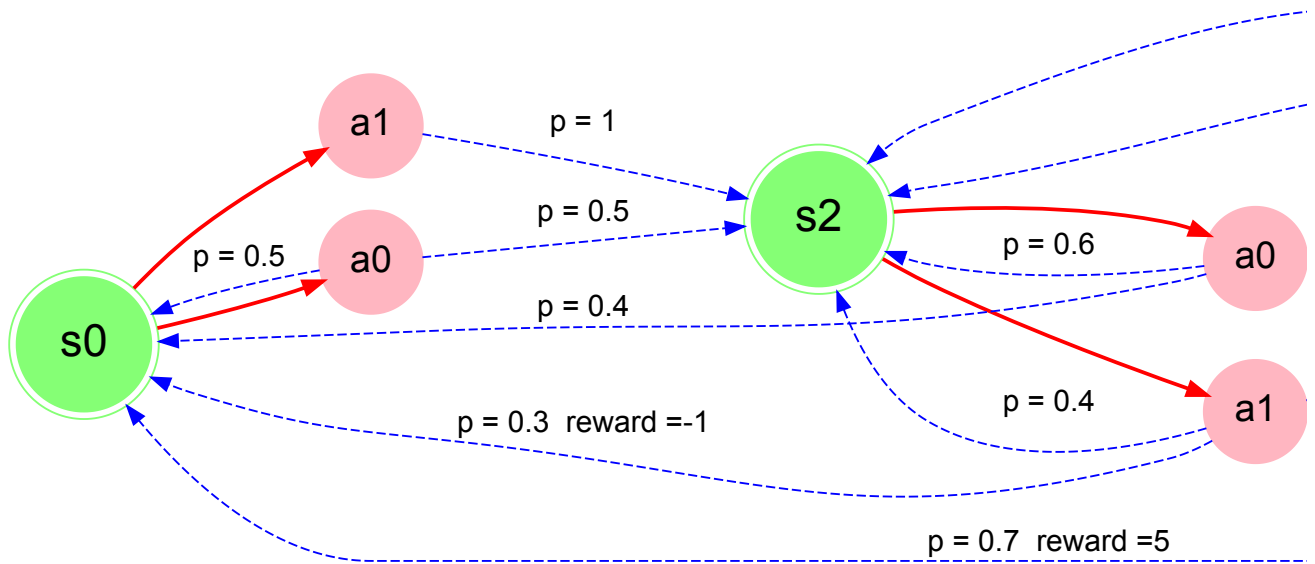
1  if has_graphviz:

```

```

2 from _mdp import plot_graph, plot_graph_with_state_values, plot_graph_opti
3 display(plot_graph(mdp))

```



## ▼ Policy iteration algorithm

Here's the pseudo-code for PI:

1. Pick an random policy  $\pi$
2. Iterate until  $\pi$  is unchanged (converges).

**Policy evaluation:**

$$V_{i+1}(s) = \sum_{s'} P(s'|s, \pi(s)) \cdot [r(s, \pi(s), s') + \gamma V_i(s')]$$

until value converges

**Policy improvement:**

$$\pi(s) = \max_a \sum_{s'} P(s'|s, a) \cdot [r(s, a, s') + \gamma V_i(s')]$$

First, let's write a function to compute the state-action value function  $Q(\pi)$ , defined as follows

$$Q_i(s, a) = \sum_{s'} P(s'|s, a) \cdot [r(s, a, s') + \gamma V_i(s')]$$

Using  $Q(s, a)$  we calculate the value at state  $s$  with policy  $\pi(s)$ .

$$V_{i+1}(s) = \sum_{s'} P(s'|s, \pi(s)) \cdot [r(s, \pi(s), s') + \gamma V_i(s')]$$

and loop until the value if coverage, the different from current loop with previous is not much.

Then, We try update the policy base on the new value of  $V$

$$\pi(s) = \max_a \sum_{s'} P(s'|s, a) \cdot [r(s, a, s') + \gamma V_i(s')]$$

loop until  $\pi$  is converges.

```

1  # Compute the Q-value using the formula above
2  def get_action_value(mdp, state_values, state, action, gamma):
3      # Initialise Q
4      Q = 0
5      for s in mdp.get_all_states():
6          # Compute Q using the equation above
7          Q = Q + mdp.get_transition_prob(state, action, s)*(mdp.get_reward(state,
8                                                                gamma*state_val
9      return Q

1  # initial simple policy
2  from collections import defaultdict
3  def init_policy():
4      policy = defaultdict(lambda: {})
5      for state in mdp.get_all_states():
6          actions = mdp.get_possible_actions(state)
7          for action in actions:
8              policy[state][action] = 1. / len(actions)
9      return policy

1  def policy_eval(mdp, policy, gamma, theta=0.00001):
2      num_iter = 100 # maximum iterations, excluding initialization
3      # initialize V(s)
4      V = {s: 0 for s in mdp.get_all_states()}
5
6      for i in range(num_iter):
7          delta = 0
8          # Compute new state values using the functions defined above.
9          new_V = {}
10         for s in mdp.get_all_states():
11             nsv = get_action_table(mdp, V, s, gamma)
12             # new_V[s] = max(nsv.items(), key=operator.itemgetter(1))[1]
13             action_policy = best_action(policy[s])
14             new_V[s] = nsv[action_policy]
15
16             delta = max(delta, abs(new_V[s] - V[s]))
17         assert isinstance(new_V, dict)
18         V = new_V
19         if delta < theta:
20             break
21
22     return V

1  def get_action_table(mdp, state_values, state, gamma):
2      A = {a: 0 for a in mdp.get_possible_actions(state)}
3      i = 0
4      # Compute all possible options
5      for a in mdp.get_possible_actions(state):
6          v = get_action_value(mdp, state_values, state, a, gamma)
7          A[a] = v
8      return A

```

```

1 def best_action(policy_state):
2     return max(policy_state.items(), key=operator.itemgetter(1))[0] # argmax

1 import operator
2 def policy_improvement(mdp, policy_eval_fn=policy_eval, gamma=0.9):
3     num_iter = 100 # maximum iterations, excluding initialization
4
5     policy = init_policy()
6     for i in range(num_iter):
7         V = policy_eval_fn(mdp, policy, gamma)
8         print(V)
9         # Will be set to false if we make any changes to the policy
10        policy_stable = True
11        # For each state...
12        for state in mdp.get_all_states():
13            # The best action we would take under the current policy
14            chosen_a = best_action(policy[state])
15
16            action_values = get_action_table(mdp, V, state, gamma)
17            best_a = best_action(action_values)
18            # Greedily update the policy
19            if chosen_a != best_a:
20                policy_stable = False
21            for action in policy[state]:
22                if action == best_a:
23                    policy[state][action] = 1
24                else:
25                    policy[state][action] = 0
26        # If the policy is stable we've found an optimal policy. Return it
27        if policy_stable:
28            return policy, V
29    return policy, V

1 import pprint
2 policy, v = policy_improvement(mdp)
3 print("Policy Probability Distribution:")
4 pprint.pprint(policy)
5 print("")
6 print(v)

{'s0': 0.0, 's1': 3.8461536621935, 's2': 0.0}
{'s0': 2.8398740312453152, 's1': 6.498777111635999, 's2': 3.4709784640026404}
{'s0': 3.7898252943341277, 's1': 7.302796844653717, 's2': 4.210930696013536}
Policy Probability Distribution:
defaultdict(<function init_policy.<locals>.<lambda> at 0x7f651d5efb90>,
          {'s0': {'a0': 0, 'a1': 1},
           's1': {'a0': 1, 'a1': 0},
           's2': {'a0': 0, 'a1': 1}})

{'s0': 3.7898252943341277, 's1': 7.302796844653717, 's2': 4.210930696013536}

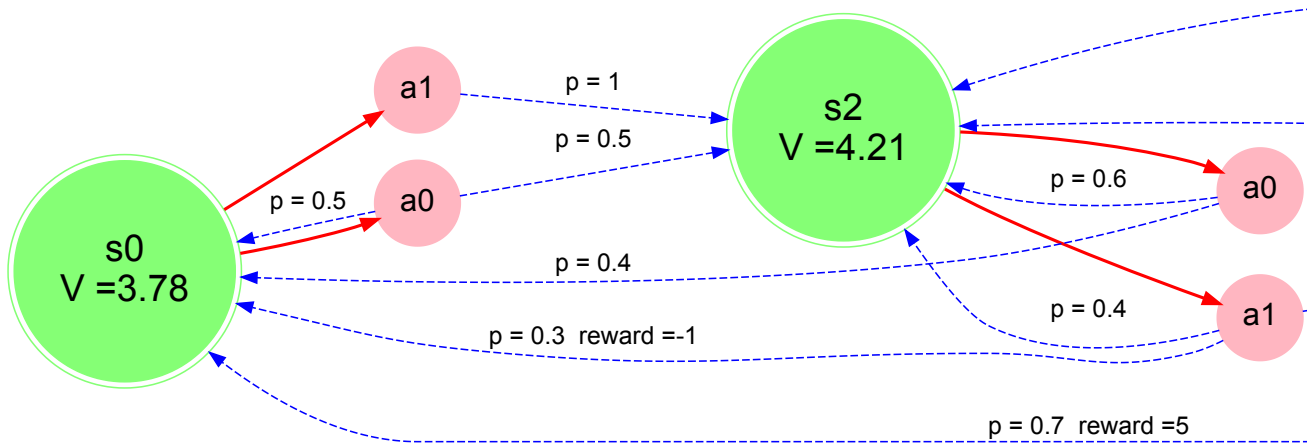
```

Finally, we can now plot the final graphical model

```

1 if has_graphviz:
2     display(plot_graph_with_state_values(mdp, v))

```

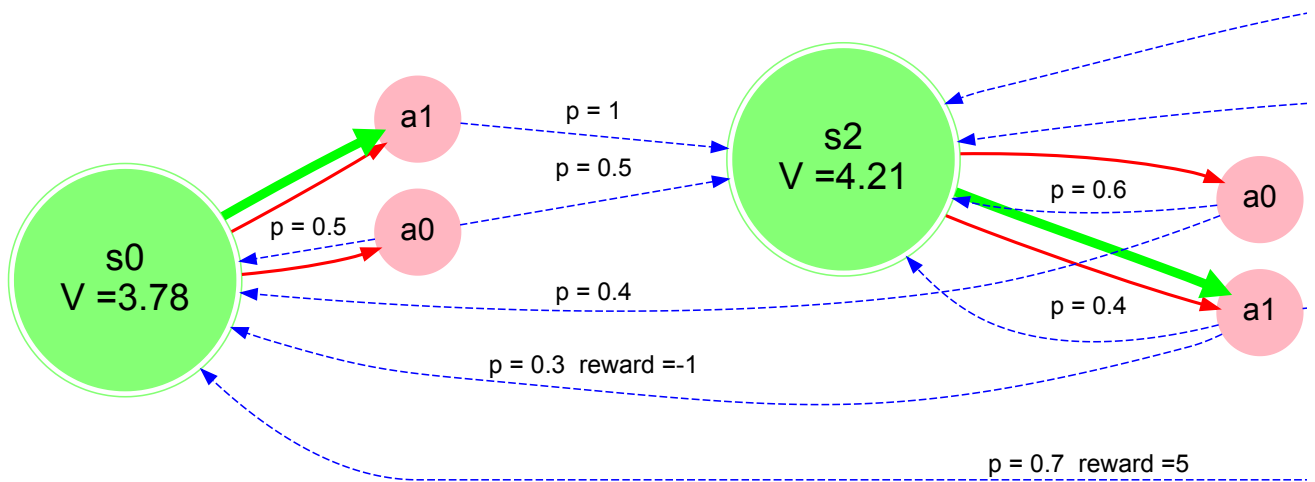


Also, we can now look at the final state values.

```

1 if has_graphviz:
2     display(plot_graph_optimal_strategy_and_state_values(mdp, v, get_action_val

```



```

1 def get_optimal_action(mdp, state):
2     # Finds optimal action using formula above.
3     if mdp.is_terminal(state):
4         return None
5
6     return best_action(policy[state])

```

```

1 gamma=0.9
2 assert get_optimal_action(mdp, 's0') == 'a1'
3 assert get_optimal_action(mdp, 's1') == 'a0'
4 assert get_optimal_action(mdp, 's2') == 'a1'

```

```

1 import numpy as np

```

```
2 # Measure agent's average reward
3
4 s = mdp.reset()
5 rewards = []
6 gamma=0.9
7
8 for _ in range(1000):
9     s, r, done, _ = mdp.step(get_optimal_action(mdp, s))
10    rewards.append(r)
11
12 print("average reward: ", np.mean(rewards))
13
14 assert(0.40 < np.mean(rewards) < 0.55)

average reward:  0.453
```

1

